

Liquibase

Database change management

Sergii Fesenko

Software development

- We do a lot of changes during development
- We use version control system for tracking code changes
- However, tracking DB changes is not as easy as tracking code changes

Database development

There are many different artifacts in database:

- tables
- data, that lives in the tables
- stored procedures
- views
- triggers
- etc (DBMS specific artifacts)

What is wrong with DB?

- database has persistent state (that's why we use it)
- code and database must corresponds
- each developer has local environment
- we change code very often (we agile :)

Actually, these issues are Agile specific, in case of waterfall they disappears :)

What is solution?

Store DB changes in text files (DDL and DML)
and apply them manually to DB.

A lot of projects use this approach

Manual tracking DB changes

Possible issues:

- easy to lost sync between code and DB state
- hard to recover from error during development (ex: in case of applying wrong statement to DB)
- often require to re-create DB from scratch during development
- hard to find out state of particular DB environment
- may require to support 2 versions of changes: full and incremental

What tools are available?

- [Flyway](#)
- [Liquibase](#)
- [c5-db-migration](#)
- [dbdeploy](#)
- [mybatis](#)
- [MIGRATEdb](#)
- [migrate4j](#)
- [dbmaintain](#)
- [AutoPatch](#)

What makes liquibase different?

- DDL abstraction DSL
- Support for different changelog format
 - build-in: XML, YAML, JSON and SQL
 - community-managed: groovy and closure
- Over 30 built-in [database refactorings](#)
- Rollback database changes feature
- Database diff report
- Extensibility

Liquibase major concepts

- Changelog file
- Changeset
- Changes
- Preconditions
- Contexts

Liquibase changelog file

The root of all Liquibase changes is the databaseChangeLog file.

```
<databaseChangeLog>
  <changeSet id="1" ...>
    ...
  </changeSet>
  <changeSet id="2" ...>
    ...
  </changeSet>
</databaseChangeLog>
```

Liquibase changelog file

It is possible to break up changelogs into multiple manageable pieces

```
<databaseChangeLog>
  <include file="src/db/tables.xml"/>

  <include file="src/db/views.xml"/>

  <include file="src/db/data.xml"/>
  ...
</databaseChangeLog>
```

Liquibase changeset

Changeset is a group of changes

Liquibase **attempts** to execute each changeSet in a transaction that is committed at the end, or rolled back if there is an error.

Liquibase changeset

```
<databaseChangeLog>
```

```
<!-- id, author and current file name must be unique -->
```

```
  <changeSet id="1" author="sfesenko" >
```

```
    <comment>Sample of changeset</comment> <!-- comment is optional -->
```

```
      <change .. /> <!-- will be explained later -->
```

```
<!-- rollback is a changeset that must be executed in order to “cancel” current changeset -->
```

```
  <rollback>
```

```
    <change .. />
```

```
  </rollback>
```

```
</changeSet>
```

```
</databaseChangeLog>
```

Liquibase changeset

```
<changeSet id="1" author="sfesenko" runAlways="true"> <!-- false by default -->
```

```
...
```

```
</changeset>
```

```
<changeSet id="2" author="sfesenko" runOnChange="true"> <!-- false by default -->
```

```
...
```

```
</changeset>
```

Useful for views, triggers, stored procedures

```
<changeSet id="3" author="sfesenko" failOnError="false"> <!-- true by default -->
```

```
...
```

```
</changeset>
```

Liquibase changeset preconditions

Both changelog and changeset may have **precondition** - assertion, that will be evaluated before execution of changeset.

Available preconditions:

- check for dbms type
- check for current user name
- check if changeset has been executed
- check if table exists
- check if table has column
- check if view exists
- check if FK constraint exists
- check if index exists
- check if sequence exists
- check if table has primary key (or specified primary key exists)
- arbitrary sql check

Liquibase changeset preconditions

Sample of precondition usage:

```
<changeSet id="1" author="sfesenko" >
  <preConditions onFail="MARK_RAN">
    <tableExists tableName="TEST" />
  </preConditions>
  <<dropTable cascadeConstraints="true" tableName="TEST" />
</changeSet>
```


Liquibase changeset context

It's possible to specify for each changeset in what context it should be run.

Context value can be specified on liquibase run.

```
<changeSet id="1" author="sfesenko" context="clear">  
    <delete tableName="USERS" />  
</changeset>
```

Liquibase change

Each changeset contains change which describes required modification of database.

It's possible to use DSL (recommended, dbms independent) or SQL for defining a change.

Liquibase change

```
<changeSet id="1" author="sfesenko">
  <createTable tableName="CURRENCY">
    <column name="CURRENCY" type="varchar(12)" >
      <constraints primaryKey="true"
primaryKeyName="PK_CURRENCY_CURRENCY"/>
    </column>
    <column name="ISO" type="java.lang.Integer" />
    <column name="DECIMAL_RULES" type="varchar(12)" />
  </createTable>

  <sql>
CREATE TABLE CURRENCY (
                CURRENCY          VARCHAR2(12) NOT NULL,
                ISO                NUMBER(10),
                DECIMAL_RULES     VARCHAR2(12),
                CONSTRAINT PK_CURRENCY_CURRENCY PRIMARY KEY (CURRENCY)
                );
  </sql>
</changeSet>
```

Liquibase changes/refactorings

Structural Refactorings

- Add Column
- Rename Column
- Modify Column
- Drop Column
- Alter Sequence
- Create Table
- Rename Table
- Drop Table
- Create View
- Rename View
- Drop View
- Merge Columns
- Create Stored Procedure

Data Quality Refactorings

- Add Lookup Table
- Add Not-Null Constraint
- Remove Not-Null Constraint
- Add Unique Constraint
- Drop Unique Constraint
- Create Sequence
- Drop Sequence
- Add Auto-Increment
- Add Default Value
- Drop Default Value

Referential Integrity Refactorings

- Add Foreign Key Constraint
- Drop Foreign Key Constraint
- Drop All Foreign Key Constraints
- Add Primary Key Constraint
- Drop Primary Key Constraint

Architectural Refactorings

- Create Index
- Drop Index
- Custom Refactorings

Non-Refactoring Transformations

- Insert Data
- Load Data
- Load Update Data
- Update Data
- Delete Data

Modifying Generated SQL

- Custom SQL
- Custom SQL File
- Custom Refactoring Class
- Execute Shell Command

How to run liquibase

On Demand

[Command Line](#)

[Ant](#)

[Maven](#)

Automated

[Servlet Listener](#)

[Spring Listener](#)

[JEE CDI Listener](#)

Update DB on application start

Java APIs

Liquibase can easily be embedded and executed through its Java APIs.

Liquibase command line

1. Create property file liquibase.property with connection parameters:

```
driver: oracle.jdbc.OracleDriver  
classpath: ojdbc14-10.2.0.3.0.jar  
url: jdbc:oracle:thin:@localhost:1521:oracle  
username: SQ_BANK  
password: SQ_BANK
```

2. Run liquibase as
./liquibase <command>

Liquibase quickstart

1. Generate changelog from existing database

Run `./liquibase --changeLogFile=changelog.xml generateChangeLog`

1. Dump data to xml file

Run `./liquibase --changeLogFile=data.xml --diffTypes=data generateChangeLog`

Be ready for `OutOfMemoryError`

Liquibase commands


- update
- updateSQL
- validate
- status
- dropAll
- rollbackCount
- generateChangeLog
- diff
- changelogSync
 - changelogSyncSQL
- clearCheckSums

Liquibase tips and tricks

- Use one file per changeset (greatly simplified merges)
- Use convention for file names
- Use “run on change” attribute for stored procedures, views, triggers, etc
- Decide early if rollback feature will be used
- Test both incremental and full update
- Do not change “wrong” changeset - just add new one with fix
- Consider possibility of “compressing” changesets when full update became slow


When it is good

Best case



When it is not so good

Worst case



Q & A